

PRACTICAL 2

Treadmill Garbage Collector

Matric. No.

140021043

26 November 2014

Introduction

As instructed by the practical specifications, this is a treadmill garbage collector modelled on Baker's implementation. The following are some notes about this particular GC implementation.

The “memory” is represented by a doubly-linked list. Each node in the list has an address, type, two values, color, next, and prev. The next and prev are of course pointers to the next and previous nodes in the linked list. The rest of the node attributes relate to the heap objects. In a real garbage collector, the heap object and its attributes would not be a part of the memory itself, and I could have moved all of them to a separate data structure. However, this would have just added a level of indirection to the node structure, and does not have any real benefits in this implementation. As such, all information about the heap objects is stored in the node struct.

Upon initialization, an empty linked list is created with a specified number of nodes – throughout writing and testing, I primarily used a linked list of 15 nodes. Heap objects, specified in main, are then added to the linked list. A word about the heap objects – all of the functionality laid out in the CF object representation is not fully implemented. Here is what is implemented:

- INT i
- BOOL b
- IND p
- WEAK p
- SOFT p

The following CF objects are either not implemented, or implemented with some restrictions:

- FLOAT: it's implemented, in a sense. The val1 parameter of a FLOAT node corresponds to the digits of the float to the left of the decimal, and the val2 parameter corresponds to the digits to the right. For example, the float 2.34 would have node->val1 = 2 and node->val2 = 34.
- TUP: it works the same as CONS in previous examples, and is limited to two pointers. For example, TUP 3 6 points to addresses 3 and 6.
- VECTOR: same as TUP
- CALL: can only have one argument. For example, CALL 3 6 points to object at 3 representing a function, and object at 6 representing the function argument.

Despite the limitations of some of the CF objects, you will find that the program runs as a proper treadmill garbage collector. Heap objects begin in the list as white, unchecked objects. The roots are then moved to grey, checked but not scanned. The roots are then scanned, in which they are moved to black and parsed. Any objects pointed to by the roots are moved to grey. Once all the roots have been scanned, all grey objects are scanned until there are no more grey objects (only black and white). Once this point has been reached, the marking is done. All white objects are released, and all black objects are switched to white. New heap objects are added, and the process begins again.

Test Cases

The following test cases were done with a doubly-linked list of 15 nodes. Ecru is free, white is unchecked, grey is visited but not scanned, black is live.

Test 1 – Single GC collection (no mutation)

Initial Heap (roots 21,0,3,12)

0		SOFT		21		0		white
3		VAR		4		0		white
6		IND		0		0		white
9		TUP		18		21		white
12		TUP		6		24		white
15		CALL		9		3		white
18		SOFT		0		0		white
21		TUP		9		0		white
24		WEAK		27		0		white
27		NULL		0		0		white
30		IND		15		0		white
33		(null)		0		0		ecru
36		(null)		0		0		ecru
39		(null)		0		0		ecru
42		(null)		0		0		ecru

Before collecting trash:

0		SOFT		21		0		black
3		VAR		4		0		black
6		IND		0		0		black
12		TUP		6		24		black
18		SOFT		0		0		black
9		TUP		18		21		black
21		TUP		9		0		black
15		CALL		9		3		white
24		WEAK		27		0		white
27		NULL		0		0		white

30		IND		15		0		white
33		(null)		0		0		ecru
36		(null)		0		0		ecru
39		(null)		0		0		ecru
42		(null)		0		0		ecru

After Collection trash:

Removing 15 from heap

Removing 24 from heap

Removing 27 from heap

Removing 30 from heap

0		SOFT		21		0		white
3		VAR		4		0		white
6		IND		0		0		white
12		TUP		6		9		white
18		SOFT		0		0		white
9		TUP		18		21		white
21		TUP		9		0		white
15		(null)		0		0		ecru
24		(null)		0		0		ecru
27		(null)		0		0		ecru
30		(null)		0		0		ecru
33		(null)		0		0		ecru
36		(null)		0		0		ecru
39		(null)		0		0		ecru
42		(null)		0		0		ecru

Test Case 2 - Two Rounds of GC (one mutation)

Initial Heap (roots 21,0,3,12)

0		SOFT		21		0		white
3		VAR		4		0		white
6		IND		0		0		white
9		TUP		18		21		white
12		TUP		6		24		white
15		CALL		9		3		white
18		SOFT		0		0		white
21		TUP		9		0		white
24		WEAK		27		0		white
27		NULL		0		0		white
30		IND		15		0		white
33		(null)		0		0		ecru
36		(null)		0		0		ecru
39		(null)		0		0		ecru
42		(null)		0		0		ecru

Objects added in mutation (roots 6,18)

15		INT		4		0		white
24		WEAK		3		0		white
27		TUP		6		24		white

After mutation and trash collecting:

Removing 3 from heap

Removing 12 from heap

Removing 15 from heap

Removing 24 from heap

Removing 27 from heap

0		SOFT		21		0		white
6		IND		0		0		white
3		(null)		0		0		ecru
12		(null)		0		0		ecru
15		(null)		0		0		ecru
24		(null)		0		0		ecru
27		(null)		0		0		ecru
30		(null)		0		0		ecru
33		(null)		0		0		ecru
36		(null)		0		0		ecru
39		(null)		0		0		ecru
42		(null)		0		0		ecru
18		SOFT		0		0		white
9		TUP		18		21		white
21		TUP		9		0		white

Test Case 3 - Four Rounds of GC (three mutations), while running out of memory space on the fourth mutation

Initial Heap (roots 21,0,3,12)

0		SOFT		21		0		white
3		VAR		4		0		white
6		IND		0		0		white
9		TUP		18		21		white
12		TUP		6		24		white
15		CALL		9		3		white
18		SOFT		0		0		white
21		TUP		9		0		white
24		WEAK		27		0		white
27		NULL		0		0		white
30		IND		15		0		white
33		(null)		0		0		ecru
36		(null)		0		0		ecru
39		(null)		0		0		ecru
42		(null)		0		0		ecru

Objects added in first mutation (roots 6,18)

15		INT		4		0		white
24		WEAK		3		0		white
27		TUP		6		24		white

After first mutation and trash collecting:

Removing 3 from heap

Removing 12 from heap

Removing 15 from heap

Removing 24 from heap

Removing 27 from heap

0		SOFT		21		0		white
6		IND		0		0		white
3		(null)		0		0		ecru
12		(null)		0		0		ecru
15		(null)		0		0		ecru
24		(null)		0		0		ecru
27		(null)		0		0		ecru
30		(null)		0		0		ecru
33		(null)		0		0		ecru
36		(null)		0		0		ecru
39		(null)		0		0		ecru
42		(null)		0		0		ecru
18		SOFT		0		0		white
9		TUP		18		21		white
21		TUP		9		0		white

Objects added in second mutation (roots 15)

3		IND		24		0		white
12		WEAK		18		0		white
15		TUP		3		12		white

After second mutation and trash collecting:

Removing 6 from heap

Removing 15 from heap

0		SOFT		21		0		white
18		SOFT		0		0		white
12		WEAK		18		0		white
3		TUP		9		12		white
6		(null)		0		0		ecru
15		(null)		0		0		ecru
24		(null)		0		0		ecru
27		(null)		0		0		ecru
30		(null)		0		0		ecru
33		(null)		0		0		ecru
36		(null)		0		0		ecru
39		(null)		0		0		ecru
42		(null)		0		0		ecru

9	TUP	18	21	white
21	TUP	9	0	white

Objects added in third mutation (roots 27)

3	TUP	9	12	white
6	BOOL	1	0	white
15	IND	18	0	white
24	TUP	3	12	white
27	IND	0	0	white
30	IND	0	0	white
33	TUP	12	12	white
36	NULL	0	0	white
39	CALL	3	15	white
42	INT	13	0	white
45	CALL	36	12	white

Not enough space! Only 9 spaces free

As a result, the last object (address 45) is left out, since there is no space for it in memory. Another approach could have been to expand the memory size as needed.

After third mutation and trash collecting:

Removing 12 from heap

Removing 3 from heap

Removing 6 from heap

Removing 15 from heap

Removing 24 from heap

Removing 30 from heap

Removing 33 from heap

Removing 36 from heap

Removing 39 from heap

Removing 42 from heap

0	SOFT	21	0	white
27	IND	0	0	white
12	(null)	0	0	ecru
3	(null)	0	0	ecru
6	(null)	0	0	ecru
15	(null)	0	0	ecru
24	(null)	0	0	ecru
30	(null)	0	0	ecru
33	(null)	0	0	ecru
36	(null)	0	0	ecru
39	(null)	0	0	ecru
42	(null)	0	0	ecru
18	SOFT	0	0	white
9	TUP	18	21	white
21	TUP	9	0	white

Timing Tests & Analysis

Test 1 - 1,000,000 iterations of the following heap (root 6)

0		SOFT		3		0		white
3		VAR		4		0		white
6		IND		0		0		white
9		TUP		18		21		white
12		(null)		0		0		ecru
15		(null)		0		0		ecru
18		(null)		0		0		ecru
21		(null)		0		0		ecru
24		(null)		0		0		ecru
27		(null)		0		0		ecru
30		(null)		0		0		ecru
33		(null)		0		0		ecru
36		(null)		0		0		ecru
39		(null)		0		0		ecru
42		(null)		0		0		ecru

Average time: 1 second

Test 2 - How does more roots affect runtime? (roots 6,3)

0		SOFT		3		0		white
3		VAR		4		0		white
6		IND		0		0		white
9		TUP		3		0		white
12		(null)		0		0		ecru
15		(null)		0		0		ecru
18		(null)		0		0		ecru
21		(null)		0		0		ecru
24		(null)		0		0		ecru
27		(null)		0		0		ecru
30		(null)		0		0		ecru
33		(null)		0		0		ecru
36		(null)		0		0		ecru
39		(null)		0		0		ecru
42		(null)		0		0		ecru

Average time: 1

Test 3 - How does number of objects affect runtime?

0		SOFT		21		0		white
3		VAR		4		0		white
6		IND		0		0		white
9		TUP		18		21		white
12		TUP		6		24		white
15		CALL		9		3		white
18		SOFT		0		0		white
21		TUP		9		0		white
24		WEAK		27		0		white
27		NULL		0		0		white
30		IND		15		0		white
33		(null)		0		0		ecru
36		(null)		0		0		ecru
39		(null)		0		0		ecru
42		(null)		0		0		ecru

Average time: 1.4 seconds

Test 4 - How does size of linked list affect runtime?

0		SOFT		21		0		white
3		VAR		4		0		white
6		IND		0		0		white
9		TUP		18		21		white
12		TUP		6		24		white
15		CALL		9		3		white
18		SOFT		0		0		white
21		TUP		9		0		white
24		WEAK		27		0		white
27		NULL		0		0		white
30		IND		15		0		white
33		(null)		0		0		ecru
36		(null)		0		0		ecru
39		(null)		0		0		ecru
42		(null)		0		0		ecru
...	
(n-1)*3		(null)		0		0		ecru

Average time for:

10 nodes: 1.7 seconds

50 nodes: 5.8 seconds

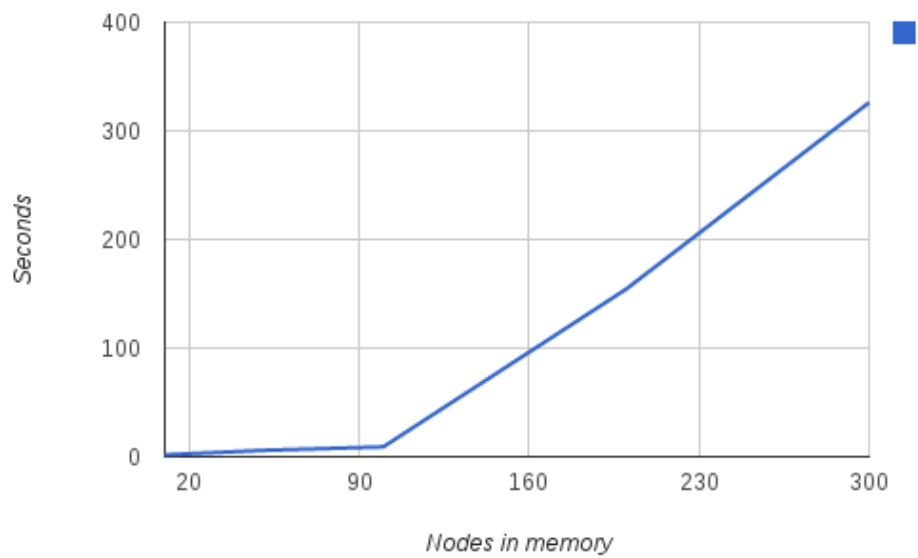
75 nodes: 8 seconds

100 nodes: 9.1 seconds

200 nodes: 154.33 seconds

500 nodes: computer crashed


Running for 1,000,000 iterations



Average time for 100,000 iterations

150 nodes: 1.7 seconds
300 nodes: 3.1 seconds
450 nodes: 4.2 seconds
1000 nodes: 13.6 seconds
2000 nodes: 515 seconds
5000 nodes: timed out

Running for 100,000 iterations



As you can see, the number of nodes in the memory linked list has an exponential effect on the runtime of the garbage collection. This is mainly due to inefficient methods in my program. Many programs, such as `find_node()`, iterate through the list to find the node it's looking for, and `find_node()` is called very frequently. Methods such as `take_out_trash()` and `reset_heap()` also contain for loops that iterate through the entire `HEAP_SIZE` variable (which I now realize is misnamed and should be called `MEMORY_SIZE`), which can be thousands of nodes for some of the higher memory allocations.